

Collaboration and Competition Project Report

Implementation

The agents are trained using a modified **Deep Deterministic Policy Gradient (DDPG)** algorithm to solve the multi-agent **Tennis** environment with two agents.


DDPG Model Architecture

The DDPG agent consists of two deep Neural Network:

- 1) Actor
- 2) Critic

The **Actor** is used to *approximate* the **optimal policy** deterministically, by learning $\operatorname{argmax}_a Q(s, a)$, which is the best action.

Then the **Critic** learns to evaluate the **optimal value function** by using the Actor's best-believed action.




Both networks (Actor and Critic) consist of **three fully connected layers**.

The first hidden layer takes the size of the state space ($state_size = 8$) as input, and outputs **400 units**, which are passed as input to the second hidden layer.

The second hidden layer outputs *300 units*, which are passed as input to the third and final output layer.

In the **Critic**, the action size is added to the input of the second hidden layer, and the final layer outputs a single Q value.

In the **Actor**, the final output from the third output layer is the size of the action space (*action_size = 2*), corresponding to the movement toward (or away from) the net, and jumping.



Batch Normalization is applied to the output of the **first fully connected** layer to normalize each dimension across the samples in a mini-batch to have unit mean and variance and to minimize covariance shift during training.

The output of the first and second fully connected layers passes through a **ReLU activation function**.

The final output layer of the **Actor** is a **tanh layer**, to bind the actions.

➡ A single Actor and Critic are used for both agents.

Learning Algorithm

- Randomly initialize critic network (*critic_local*) and actor network (*actor_local*) with random weights θ^Q and θ^π .
- Initialize target networks (*actor_target*) and (*critic_target*) with weights $\theta^{Q'} \leftarrow \theta^Q$, $\theta^{\pi'} \leftarrow \theta^\pi$.
- Initialize replay memory (*ReplayBuffer*) with capacity (*BUFFER_SIZE = 10⁶*).

- **for** episode ($i_episode \leftarrow 1$ to 1000):
 - Prepare initial state: ($states = env_info.vector_observations$)
 - Initialize $scores \leftarrow 0$
 - **for** time step $t \leftarrow 1$ to 1000 :
 - Select *actions* A from *states* S according to the current policy and exploration noise (an Ornstein-Uhlenbeck process with $\theta = 0.15$ and $\sigma = 0.2$).
 - Execute *actions* A, observe *rewards* R
 - Prepare next state: ($next_states = env_info.vector_observations$)
 - Store experience tuple (S,A,R,S') in replay memory (*ReplayBuffer*)
 - $states \leftarrow next_states$
 - Add reward to score
- Every ($LEARN_EVERY = 1$) timesteps:
- Sample a random mini-batch of experience tuples (states, actions, rewards, next_states, dones) from memory (*ReplayBuffer*)
 - Compute Q targets for current states $Q_targets = rewards + (gamma * Q_targets_next * (1 - dones))$
 - Update critic by minimizing the loss and applying Gradient Clipping to avoid exploding gradient problem.
 - Update the actor policy using the sampled policy gradient.
 - Update the target networks with ($\tau = 0.002$) of the local network weights: $\theta_target = \tau * \theta_local + (1 - \tau) * \theta_target$.

Hyperparameters

Variable Name	Chosen Value	Description
BUFFER_SIZE	10^6	Replay buffer size
BATCH_SIZE	256	Mini-batch size
GAMMA	0.99	Discount Factor
TAU	0.002	Soft target updates value
LR_ACTOR	0.001	Actor's Learning Rate
LR_CRITIC	0.001	Critic's Learning Rate
WEIGHT_DECAY	0	L2 weight decay

LEARN_EVERY	1	Learning timestep interval
LEARN_NUM	1	Number of learning passes
GRAD_CLIPPING	1	Gradient Clipping value
OU_SIGMA	0.01	Sigma for Ornstein-Uhlenbeck noise process
OU_THETA	0.15	Theta for Ornstein-Uhlenbeck noise process

Results

Episode 100	Max. Score: 0.10	Avg. Score: 0.01
Episode 200	Max. Score: 0.20	Avg. Score: 0.03
Episode 300	Max. Score: 0.90	Avg. Score: 0.07
Episode 400	Max. Score: 1.70	Avg. Score: 0.16
Episode 440	Max. Score: 2.50	Avg. Score: 0.51
Environment solved in 340 episodes!		Average Score: 0.51
Training complete in 57m 37s		




Figure 1 Plot of Rewards

The agents were able to receive an average score of **0.51** (over 100 consecutive episodes, after taking the maximum over both agents) in **340 episodes**.

Future Improvements

Several improvements could be implemented to enhance the agent's performance, including:

- Implementing [Prioritized Experience Replay](#) that replays **important** transitions *more frequently*, and therefore learns more efficiently.
- Implementing Multi-Agent DDPG (MADDPG) and compare its performance to simple DDPG.
- Implementing more stable methods to achieve better performance, like: Trust Region Policy Optimization (TRPO), Truncated Natural Policy Gradient (TNPG), Proximal Policy Optimization (PPO) or the more recent [Distributed Distributional Deterministic Policy Gradients \(D4PG\)](#).
- Fine tuning the hyperparameters further to solve the environment.

References

- 1) Continuous control with deep reinforcement learning
<https://arxiv.org/abs/1509.02971>
- 2) DDPG Pendulum Exercise: Udacity's Deep Reinforcement Learning Nanodegree
GitHub Repo
<https://github.com/udacity/deep-reinforcement-learning/tree/master/ddpg-pendulum>
- 3) Multi-Agent Actor-Critic for Mixed Cooperative-Competitive Environments
<https://arxiv.org/abs/1706.02275>
- 4) Benchmarking Deep Reinforcement Learning for Continuous Control
<https://arxiv.org/abs/1604.06778>
- 5) Prioritized Experience Replay
<https://arxiv.org/abs/1511.05952>